

# CS684 Embedded Systems (Software)

#### **Embedded Applications**

Kavi Arya Paritosh Pandya

CSE/ IIT Bombay

# Examples of Embedded Systems





We look at details of

Simple Digital Camera

# Embedded Applications



## They are everywhere!

- Wristwatches, washing machines,
- Microwave ovens,
- Elevators, mobiles, printers
- Telephone exchanges,
- Automobiles, aircrafts, ...

# Common Design Metrics



- NRE (Non-recurring engineering) cost
- Unit cost
- Size (bytes, gates)
- Performance (execution time)
- Power (more power=> more heat & less battery time)
- Flexibility (ability to change functionality)
- Time to prototype
- Time to market
- Maintainability
- Correctness
- Safety (probability that system won't cause harm)

# Embedded Apps



#### A modern home

- Has a few general purpose PCs/laptops
- But dozens of embedded systems.

## More prevalent in industrial sectors

- 10's of embedded computers in modern automobiles
- Chemical and nuclear power plants

# Embedded Applications



An embedded system typically has a digital signal processor and a variety of I/O devices connected to sensors and actuators.

Computer (controller) surrounded by other subsystems, sensors and actuators

#### **Computer -- Controller's function is:**

- Monitor parameters of physical processes of "environment"
- Control these processes whenever needed.

# Simple Examples



#### A simple thermostat controller

- Periodically reads temperature of chamber
- Switches on or off the cooling system.

#### A pacemaker

- Constantly monitors the heart
- Paces heart when heart beats are missed

# 1. Digital Camera: An Embedded System



- Introduction to simple digital camera
- Requirements specification
- Designer's perspective

Design exploration



# Requirements Specification



- System's reqmts what system should do
  - Nonfunctional requirements
    - Constraints on design metrics (e.g., "should use 0.001 watt or less")
  - Functional requirements
    - System's behavior
       (e.g., "output X to be input Y times 2")

— . . . .

# Requirements Specification...



# Initial specification may be general and come from marketing dept.

- E.g., short document detailing market need for a lowend digital camera that:
  - Captures/ stores at least 50 low-res images and uploads to PC,
  - Costs around \$100 with single medium-size IC costing < \$25,</li>
  - As long as possible battery life,
  - Expected sales volume of 200,000 if market entry < 6 months,</li>
  - 100,000 if between 6 and 12 months,
  - Insignificant sales beyond 12 months

# Nonfunctional requirements



- Design metrics of importance based on initial specification
  - Performance: time required to process image
  - Size: number of elementary logic gates (2-input NAND gate) in IC
  - Power: measure of avg. electrical energy consumed while processing
  - Energy: battery lifetime (power x time)

# Nonfunctional requirements...



- Constrained metrics
  - Values <u>must</u> be below (sometimes above) certain threshold
- Optimization metrics
  - Improve as much as possible to improve product
- Metric can be both constrained and optimization

# Nonfunctional requirements...



#### Power

- Must operate below certain temperature (cooling fan not possible)
- Therefore, constrained metric

#### Energy

- Reducing power or time reduces energy
- Optimized metric: battery to last as long as possible

# Nonfunctional requirements...



#### Performance

- Must process image fast enough to be useful
- 1 sec reasonable constraint
  - Slower would be annoying
  - Faster not necessary for low-end of market
- Therefore, constrained metric

#### Size

- Must use IC that fits in reasonably sized camera
- Constrained and optimization metric
  - Constraint may be 1M gates, but smaller would be cheaper

# Example: Panasonic Lumix DMC TZ5







\$300

- 9.1 effective Megapixels
- 28-280mm equiv lens, 10x optical zoom & 4x Digital Zoom
- 3.0-inch LCD with 460,000 dots resolution
- Optical Image Stabilizer
- ISO sensitivity up to 6400
- Face Detection AF
- 6 shooting modes, 23 scene modes inc. Intelligent Auto mode
- Venus Engine IV processor
- HD output
- In-Camera Editing

1. Digital Camera: An Embedded System

#### Design

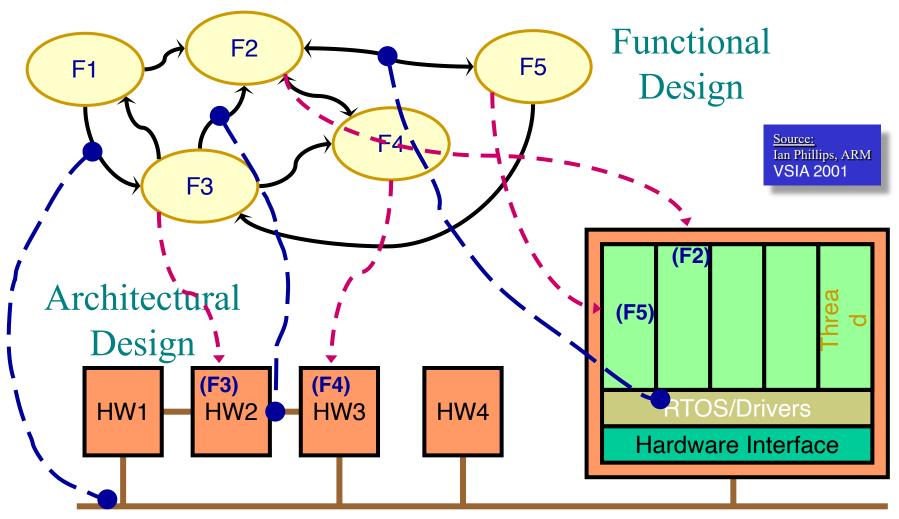
- Four implementations
- -Issues:
  - General-purpose vs. single purpose processors?
  - Partitioning of functionality among different processor types?



© Kavi Arya IIT Bombay 16

## Functional Design & Mapping





# Introduction to a simple digital camera



- Captures images
- Stores images in digital format
  - Multiple images stored in camera
    - Number depends on memory and bits/image
- Downloads images to PC
  - Serial comm (USB, etc.)
  - Wireless (Bluetooth, 802.11, ...)

© Kavi Arya IIT Bombay 18

# Introduction to a simple digital camera...



## Only possible in couple of decades

- Systems-on-a-chip
  - Multiple processors and memories on one IC
- High-capacity flash memory

## Very simple description used for example

- Many more features with real digital camera
  - Variable size images, image deletion, digital stretching, zooming in and out, etc.

© Kavi Arya IIT Bombay 19

# Designer's perspective



## Two key tasks

- 1. Processing images and storing in memory
  - When shutter pressed:
    - Image captured
    - Converted to digital form by charge-coupled device (CCD)
    - Compressed and archived in internal memory
- 2. Uploading images to PC
  - Digital camera attached to PC
  - Software to transmit archived images serially

# Charge-coupled device (CCD)

- Special sensor that captures an image
- Light-sensitive silicon solid-state device composed of many cells

#### Light

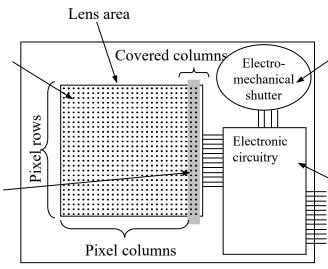
⇒charge

 $\Rightarrow$ 8-bit value:

0 => no exposure

255=> intense light

Some columns covered with black strip.
Light-intensity here used for zerobias adjustment



Electromechanical shutter activates to expose cells to light

Circuitry discharges cells, activates shutter, reads 8-bit value of each cell.

Values clocked out of CCD by external logic through std parallel bus interface.

#### Zero-bias error

- Manufacturing errors cause cells to measure slightly above or below actual light intensity
- Error typically same across columns, but different across rows
- Some of left most columns blocked by black paint to detect zero-bias error
  - Reading of non-zero in blocked cells is zero-bias error
  - Each row corrected by subtracting avg error in blocked cells for that row

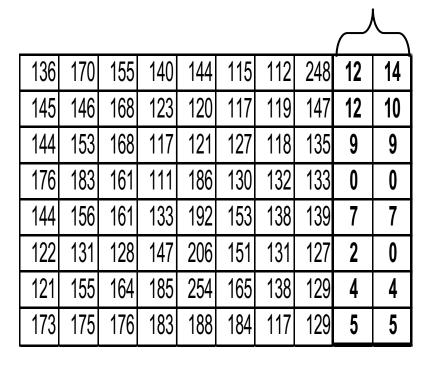
© Kavi Arya IIT Bombay 22

#### Zero-bias error...



#### **Covered cells**

#### Zero-bias adjustment



-13	
-11	
<b>-9</b>	
0	
<b>-</b> 7	
-1	
-4	
<b>-</b> 5	

123	157	142	127	131	102	99	235
134	135	157	112	109	106	108	136
135	144	159	108	112	118	109	126
176	183	161	111	186	130	132	133
137	149	154	126	185	146	131	132
121	130	127	146	205	150	130	126
117	151	160	181	250	161	134	125
168	170	171	178	183	179	112	124

Before zero-bias adjustment

After zero-bias adjustment

# Compression



- Store more images
- Transmit image to PC in less time
- JPEG (Joint Photographic Experts Group)

© Kavi Arya IIT Bombay 24

# Compression...



#### JPEG (Joint Photographic Experts Group)

- Popular standard format for representing digital images in a compressed form
- Provides for a number of different modes of operation
- Sequential Mode used here provides high compression ratios using DCT (Discrete Cosine Transform)

(others are -- progressive, lossless, hierarchical)

- Image data divided into blocks of 8 x 8 pixels
- 3 steps performed on each block

DCT, Quantization, Huffman encoding

# DCT step



- Transforms original 8 x 8 block into a cosine-frequency domain
  - Upper-left corner values represent more of essence of image

(Average for the image)

- Lower-right corner values represent finer details
  - Can reduce precision of these values and retain reasonable image quality
  - Quantize many may become 0

© Kavi Arya IIT Bombay 26

# DCT step...



## FDCT (Forward DCT) formula

- C(h) = if (h == 0) then 1/sqrt(2) else 1.0
  - Auxiliary function used in main function F(u,v)
- F(u,v) =  $\frac{1}{4}$  x C(u) x C(v) Σx=0..7 Σy=0..7 Dxy x cos(π(2u + 1)u/16) x cos(π(2y + 1)v/16)
  - Gives encoded pixel at row u, column v
  - Dxy is original pixel value at row x, column y

# IDCT (Inverse DCT)

 Reverses process to obtain original block (not needed for this design)

# Quantization step



- Achieve high compression ratio by reducing image quality
  - Reduce bit precision of encoded data
    - Fewer bits needed for encoding
    - One way is divide all values by factor of 2
      - -Simple right shifts can do this
      - -General: table driven mapping
  - Dequantization reverses process for decompression

# Quantization step...



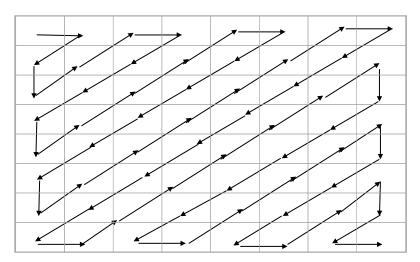
								•								
1150	39	-43	-10	26	-83	11	41		144	5	<b>-</b> 5	-1	3	-10	1	43
-81	-3	115	-73	-6	-2	22	-5		-10	0	14	<b>-</b> 9	1	0	3	~-
14	-11	1	-42	26	-3	17	-38	Divide	2	1	0	-5	3	0	2	4,7
2	-61	-13	-12	36	-23	-18	5	each cell's	0	-8	-2	-2	5	-3	-2	
44	13	37	-4	10	-21	7	-8	value by 8	6	2	5	-1	1	-3	1	1
36	-11	<b>-</b> 9	<b>-</b> 4	20	-28	-21	14		5	1	1	-1	3	<b>-</b> 4	-3	2
-19	-7	21	-6	3	3	12	-21		-2	-1	3	-1	0	0	2	-3
-5	-13	-11	-17	-4	-1	7	-4		-1	-2	-1	-2	-1	0	1	•

After being decoded using DCT

After quantization

# Huffman encoding step

- Serialize 8 x 8 block of pixels
  - Values are converted into single list using zigzag pattern



Usually, first item of blocks are stored differentially

Zigzag brings equal values together => run-length encoding

© Kavi Arya IIT Bombay 30

# Huffman encoding step...



## Perform Huffman encoding

- More frequently occurring pixels assigned short binary code
- Longer binary codes left for less frequently occurring pixels

## Each pixel in serial list converted to Huffman encoded values

Much shorter list, thus compression

# Huffman encoding example...



#### Pixel frequencies on left table:

- Pixel value –1 occurs 15 times
- Pixel value 14 occurs 1 time

#### Build Huffman tree from bot up

- Create one leaf node for each pixel value and assign frequency as node's value
- Create internal node by joining any two nodes whose sum is a minimal value. This sum is internal node's value
- Repeat until complete binary tree
- Traverse tree from root to leaf to obtain binary code for leaf's pixel value
  - Append 0 for left traversal, 1 for right traversal

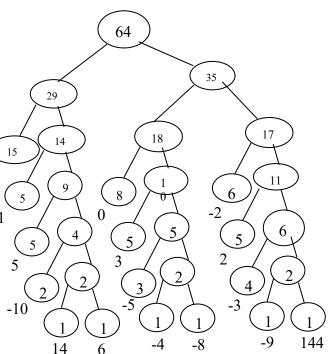
#### Huffman encoding is reversible

No code is prefix of another code

Pixel frequencies

-	4.5
-1	15x
0	8x
-2	6x
1	5x
2	5x
3	5x
5	5x
-3	4x
-5	4x 3x
-10	2x
144	1x
-9	1x
-8	1x
-4	1x
6	1x
14	1x

-1



Huffman tree

Huffman codes

-1	00					
0	100					
-2	110					
1	010					
2	1110					
3	1010					
5	0110					
-3	11110					
-5	10110					
-10	01110					
144	111111					
-9	111110					
-8	101111					
-4	101110					
6	011111					
14	011110					

# Archive step



- Record starting address and image size
  - Can use linked list
- One possible way to archive images
  - If max number of images archived is N:
    - Set aside memory for N addresses and N image-size variables
    - Keep counter for location of next available address
    - Initialize addresses and image-size variables to 0
    - Set global memory address to N x 4
      - Assuming addresses, image-size variables occupy N x 4 bytes
    - First image archived starting at address N x 4
    - Global memory address updated to N x 4 + (compressed image size)
- Memory requirement based on:
  - N, image size, and average compression ratio

# Uploading to PC

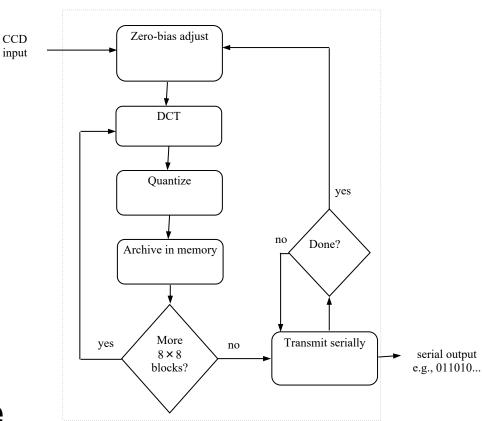


- When connected to PC and upload command received
  - Read images from memory
  - Transmit serially using UART\*
  - While transmitting
- Reset pointers, image-size variables and global memory pointer accordingly

# Informal functional specification

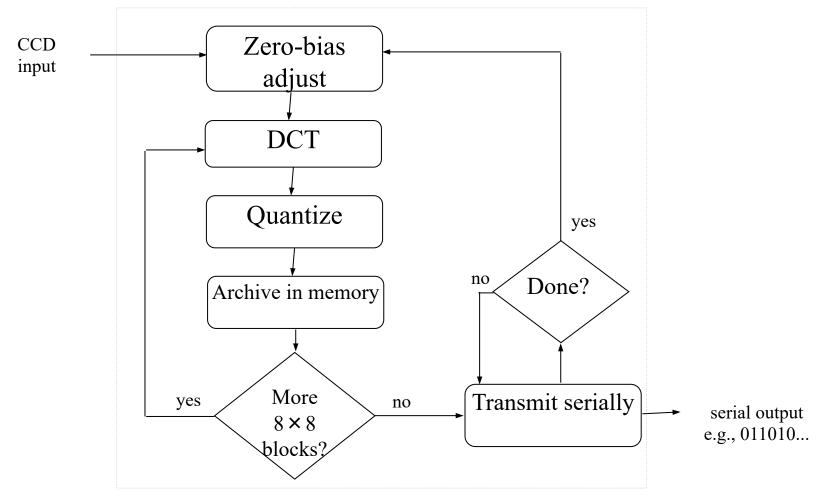


- Flowchart breaks functionality down into simpler functions
- Each function's details described in English
- Low quality image has resolution of 64 x 64
- Mapping functions to a particular processor type not done at this stage



# Informal functional specification





# Refined functional specification

 Refine informal specification into one that can actually be executed

- Can use C-like code to describe each function
  - Called system-level model, prototype, or simply model
  - Also is first implementation

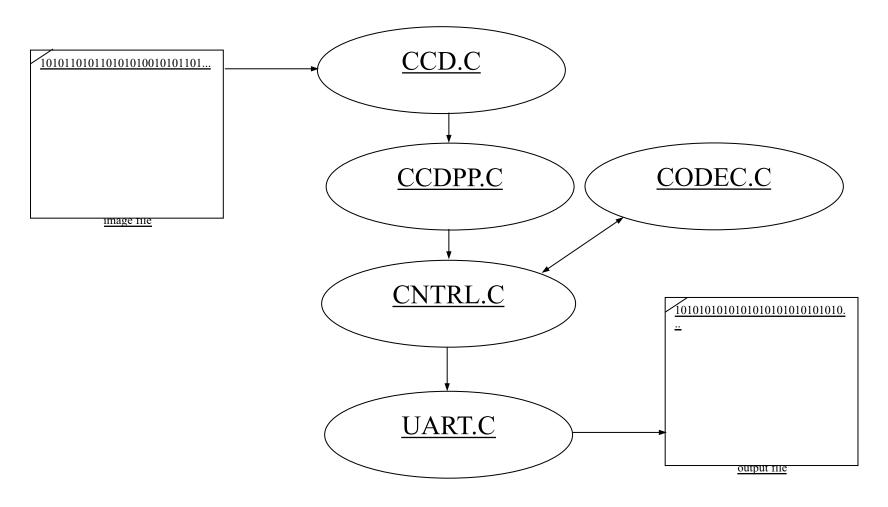
UART.C

101010...

output file

# Executable model of digital camera









- Provides insight into operations of system
  - Profiling finds computationally intensive functions
- Can obtain sample output used to verify correctness of final implementation

## CCD module



- Simulates real CCD
- CcdInitialize is passed name of image file
- CcdCapture reads "image" from file into buffer
- CcdPopPixel outputs pixels one at a time from buffer





- Performs zero-bias adjustment
- CcdppCapture uses
   CcdCapture and
   CcdPopPixel to obtain image
- Performs zero-bias adjustment after each row read in

# CCDPP (CCD PreProcessing) module



- Performs zero-bias adjustment
- CcdppCapture uses CcdCapture and CcdPopPixel to obtain image
- Performs zero-bias adjustment after each row read in

```
void CcdppCapture(void) {
    char bias;
    CcdCapture();
    for(rowIndex=0; rowIndex<SZ_ROW; rowIndex++) {
        for(colIndex=0; colIndex<SZ_COL; colIndex++) {
            buffer[rowIndex][colIndex] = CcdPopPixel();
        }
        bias = (CcdPopPixel() + CcdPopPixel()) / 2;
        for(colIndex=0; colIndex<SZ_COL; colIndex++) {
            buffer[rowIndex][colIndex] -= bias;
        }
    }
    rowIndex = 0;
    colIndex = 0;</pre>
```

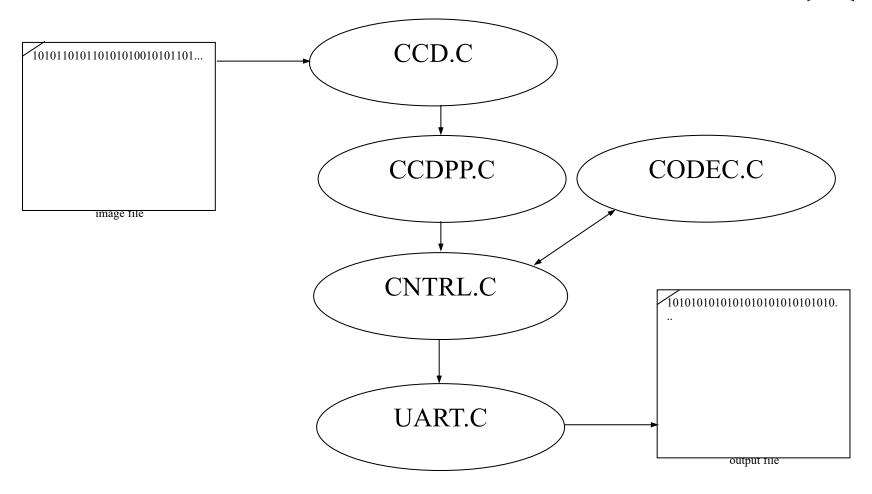
```
#define SZ_ROW 64
#define SZ_COL 64
static char buffer[SZ_ROW][SZ_COL];
static unsigned rowIndex, colIndex;
```

```
void CcdppInitialize() {
   rowIndex = -1;
   colIndex = -1;
}
```

```
char CcdppPopPixel(void) {
   char pixel;
   pixel = buffer[rowIndex][colIndex];
   if( ++colIndex == SZ_COL ) {
      colIndex = 0;
      if( ++rowIndex == SZ_ROW ) {
       colIndex = -1;
       rowIndex = -1;
    }
   }
   return pixel;
}
```

# Executable model of digital camera





#### UART module



- Actually a half UART
  - Only transmits, does not receive
- UartInitialize is passed name of file to output to
- UartSend transmits (writes to output file) bytes at a time

```
#include <stdio.h>
static FILE *outputFileHandle;
void UartInitialize(const char *outputFileName) {
    outputFileHandle = fopen(outputFileName, "w");
}
void UartSend(char d) {
    fprintf(outputFileHandle, "%i\n", (int)d);
}
```

## CODEC module



- Models FDCT\* encoding
- ibuffer holds original 8 x 8 block
- obuffer holds encoded 8 x 8 block
- CodecPushPixel called 64x to fill ibuffer w/original block
- CodecDoFdct called once to transform 8 x 8 block
  - Explained in next slide
- CodecPopPixel called 64 times to retrieve encoded block from obuffer

## CODEC module



- Models FDCT encoding
- ibuffer holds original 8 x 8 block
- obuffer holds encoded 8 x 8 block
- CodecPushPixel called 64 times to fill ibuffer with original block
- CodecDoFdct called once to transform 8 x 8 block
  - Explained in next slide
- CodecPopPixel called 64 times to retrieve encoded block from obuffer

```
static short ibuffer[8][8], obuffer[8][8], idx;

void CodecInitialize(void) { idx = 0; }
```

```
void CodecPushPixel(short p) {
   if( idx == 64 ) idx = 0;
   ibuffer[idx / 8][idx % 8] = p; idx++;
}
```

```
void CodecDoFdct(void) {
   int x, y;
   for(x=0; x<8; x++) {
      for(y=0; y<8; y++)
            obuffer[x][y] = FDCT(x, y, ibuffer);
    }
   idx = 0;
}</pre>
```

```
short CodecPopPixel(void) {
    short p;
    if( idx == 64 ) idx = 0;
    p = obuffer[idx / 8][idx % 8]; idx++;
    return p;
}
```

# FDCT (Forward DCT) formula



$$C(h) = if (h == 0) then 1/sqrt(2) else 1.0$$

Auxiliary function used in main function F(u,v)

```
F(u,v) = \frac{1}{4} \times C(u) \times C(v)
\Sigma x = 0..7 \ \Sigma y = 0..7 \ Dxy \times \cos(\pi(2x + 1)u/16) \times \cos(\pi(2y + 1)v/16)
= \frac{1}{4} \times C(u) \times C(v)
\Sigma x = 0..7 \cos(\pi(2x + 1)u/16) \times \Sigma y = 0..7 \ Dxy \times \cos(\pi(2y + 1)v/16)
```

- Gives encoded pixel at row u, column v
- Dxy is original pixel value at row x, column y

## CODEC...



- Implementing FDCT formula
- Only 64 possible inputs to COS, so table can be used to save performance time
  - Floating-point values multiplied by 32,678 and rounded to nearest integer
  - 32,678 chosen to store each value in 2 bytes of memory
  - Fixed-point representation explained more later
- FDCT unrolls inner loop of summation, implements outer summation as two consecutive for loops

#### CODEC...



- Implementing FDCT formula
- Only 64 possible inputs to COS, so table can be used to save performance time
  - Floating-point values multiplied by 32,678 and rounded to nearest integer
  - 32,678 chosen in order to store each value in 2 bytes of memory
  - Fixed-point representation explained more later
- FDCT unrolls inner loop of summation, implements outer summation as two consecutive for loops

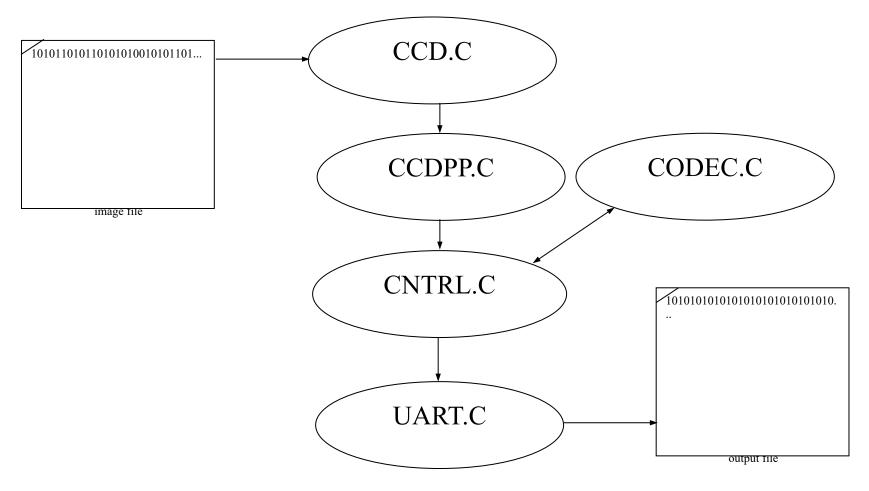
```
static short ONE_OVER_SQRT_TWO = 23170;
static double COS(int xy, int uv) {
    return COS_TABLE[xy][uv] / 32768.0;
}
static double C(int h) {
    return h ? 1.0 : ONE_OVER_SQRT_TWO / 32768.0;
}
```

```
static const short COS_TABLE[8][8] = {
      { 32768, 32138, 30273, 27245, 23170, 18204, 12539, 6392 },
      { 32768, 27245, 12539, -6392, -23170, -32138, -30273, -18204 },
      { 32768, 18204, -12539, -32138, -23170, 6392, 30273, 27245 },
      { 32768, 6392, -30273, -18204, 23170, 27245, -12539, -32138 },
      { 32768, -6392, -30273, 18204, 23170, -27245, -12539, 32138 },
      { 32768, -18204, -12539, 32138, -23170, -6392, 30273, -27245 },
      { 32768, -27245, 12539, 6392, -23170, 32138, -30273, 18204 },
      { 32768, -32138, 30273, -27245, 23170, -18204, 12539, -6392 }
};
```

```
static int FDCT(int u, int v, short img[8][8]) {
   double s[8], r = 0; int x;
   for(x=0; x<8; x++) {
      s[x] = img[x][0] * COS(0, v) + img[x][1] * COS(1, v) +
            img[x][2] * COS(2, v) + img[x][3] * COS(3, v) +
            img[x][4] * COS(4, v) + img[x][5] * COS(5, v) +
            img[x][6] * COS(6, v) + img[x][7] * COS(7, v);
   }
   for(x=0; x<8; x++) r += s[x] * COS(x, u);
   return (short)(r * .25 * C(u) * C(v));</pre>
```

# Executable model of digital camera





# CNTRL (controller) module



- Heart of the system
- CntrlCaptureImage uses CCDPP module to input image and place in buffer
- CntrlCompressImage breaks the 64 x 64 buffer into 8 x 8 blocks and performs FDCT on each block using the CODEC module
  - Also performs quantization on each block
- CntrlSendImage transmits encoded image serially using UART module

## CNTRL (controller) module



- Heart of the system
- Cntrllnitialize for consistency with other modules only
- CntrlCaptureImage uses CCDPP module to input image and place in buffer
- CntrlCompressImage breaks the 64 x 64 buffer into 8 x 8 blocks and performs FDCT on each block using the CODEC module
  - Also performs quantization on each block
- CntrlSendImage transmits encoded image serially using UART module

```
void CntrlCaptureImage(void) {
    CcdppCapture();
    for(i=0; i<SZ_ROW; i++)
        for(j=0; j<SZ_COL; j++)
        buffer[i][j] = CcdppPopPixel();
}</pre>
```

```
#define SZ_ROW 64
#define SZ_COL 64
#define NUM_ROW_BLOCKS (SZ_ROW / 8)
#define NUM_COL_BLOCKS (SZ_COL / 8)
static short buffer[SZ_ROW][SZ_COL], i, j, k, l, temp;
void CntrlInitialize(void) {}
```

# Putting it all together



- Main initializes all modules, then uses CNTRL module to capture, compress, and transmit one image
- This system-level model can be used for extensive experimentation
  - Bugs much easier to correct here rather than in later
     models [int main(int argc, char \*argv[]) {

```
int main(int argc, char *argv[]) {
    char *uartOutputFileName = argc > 1 ? argv[1] : "uart_out.txt";
    char *imageFileName = argc > 2 ? argv[2] : "image.txt";
    /* initialize the modules */
    UartInitialize(uartOutputFileName);
    CcdInitialize(imageFileName);
    CcdppInitialize();
    CodecInitialize();
    CntrlInitialize();
    /* simulate functionality */
    CntrlCaptureImage();
    CntrlSendImage();
}
```

# Design



## Determine system's architecture

- Processors
  - Any combination of single-purpose (custom or standard) or general-purpose processors
- Memories, buses

## Map functionality to that architecture

- Multiple functions on one processor
- One function on one or more processors

# Design..



#### Implementation

- A particular architecture and mapping
- Solution space is set of all implementations

#### Starting point

- Low-end gen. purpose processor connected to flash memory
  - All functionality mapped to software running on processor
  - Usually satisfies power, size, time-to-market constraints
  - If timing constraint not satisfied then try:
    - use single-purpose processors for time-critical functions
    - rewrite functional specification

# Implementation 1: Microcontroller alone

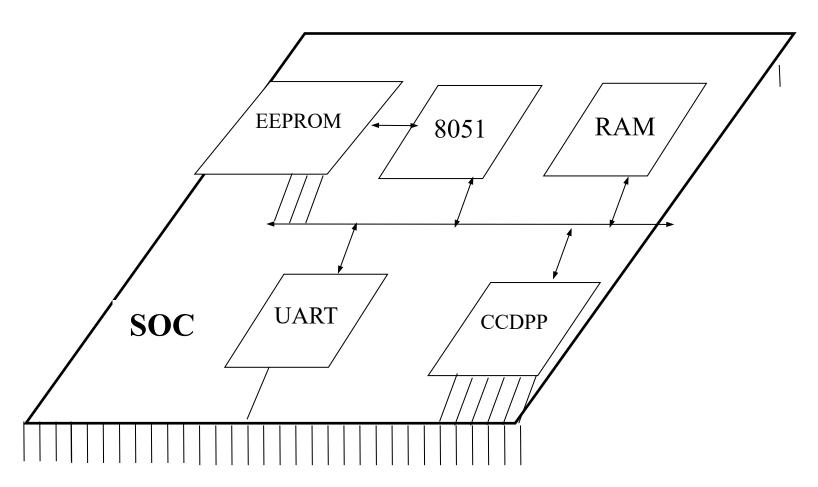
- Low-end processor could be Intel 8051 microcontroller Today: RPi, ARM Cortex,...
- Total IC cost including NRE about \$5
- Well below 200 mW power
- Time-to-market about 3 months
- However...



- However, one image per second not possible
  - 12 MHz, 12 cycles per instruction
    - Executes one million instructions per second
  - CcdppCapture has nested loops => 4096 (64x64) iterations
    - ~100 assembly instructions each iteration
    - 409,000 (4096 x 100) instructions per image
    - Half of budget for reading image alone
  - Would be over budget after adding compute-intensive DCT and Huffman encoding

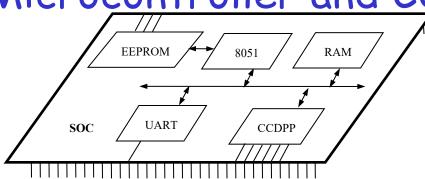
# Implementation 2: Microcontroller and CCDPP





# Implementation 2: Microcontroller and CCDPP



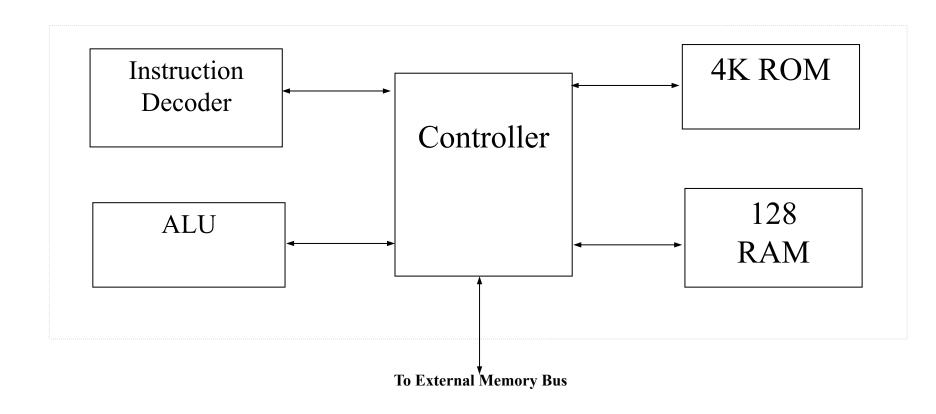


- CCDPP function on custom single-purpose processor
  - Improves performance less microcontroller cycles
  - Increases NRE cost and time-to-market
  - Easy to implement: Simple datapath, Few states in controller
- Simple UART easy to implement as single-purpose processor also
- EEPROM for program memory and RAM for data memory added as well

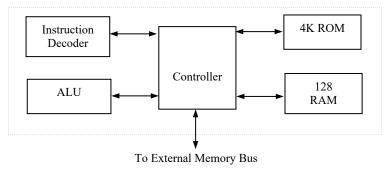
## Microcontroller



#### Block diagram of Intel 8051 processor core



#### Microcontroller



- Synthesizable version of Intel 8051 available
  - Written in VHDL
  - Captured at register transfer level (RTL)
- Fetches instruction from ROM
- Decodes using Instruction Decoder
- ALU executes arithmetic operations
  - Source and destination registers reside in RAM
- Special data movement instructions used to load and store externally
- Special program generates VHDL description of ROM from output of C compiler/linker

# Connecting SOC components



#### Memory-mapped

 All single-purpose processors and RAM are connected to 8051's memory bus

#### Read

- Processor places address on 16-bit address bus
- Asserts read control signal for 1 cycle
- Reads data from 8-bit data bus 1 cycle later
- Device (RAM or SPP) detects asserted read control signal
- Checks address
- Places and holds requested data on data bus for 1 cycle

# Connecting SOC components...



#### Write

- Processor places address/data on address/data bus
- Asserts write control signal for 1 clock cycle
- Device (RAM or SPP) detects asserted write control signal
- Checks address bus
- Reads and stores data from data bus

## Software



- System-level model provides majority of code
  - Module hierarchy, procedure names, and main program unchanged
- Code for UART and CCDPP modules must be redesigned
  - Simply replace with memory assignments
    - xdata used to load/store variables over ext. memory bus
    - \_at\_ specifies memory address to store these variables
    - Byte sent to U\_TX\_REG by processor will invoke UART
    - U\_STAT\_REG used by UART to indicate its ready for next byte
      - UART may be much slower than processor
  - Similar modification for CCDPP code
- All other modules untouched

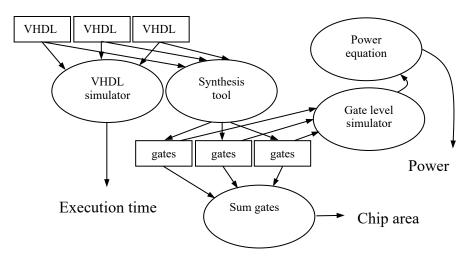
# Analysis



#### Entire SOC tested on VHDL simulator

- Interprets VHDL descriptions and functionally simulates execution of system
  - Recall program code translated to VHDL description of ROM
- Tests for correct functionality
- Measures clock cycles to process one image (performance)

#### **Obtaining design metrics of interest**

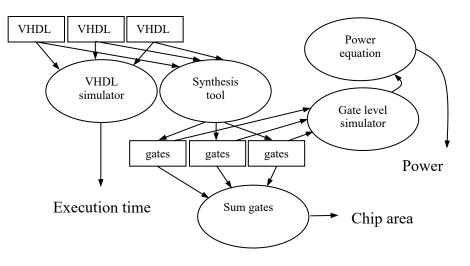


# Analysis...



- Gate-level description obtained through synthesis
  - Synthesis tool like compiler for SPPs
  - Simulate gate-level models to obtain data for power analysis
    - Number of times gates switch from: 1 to 0 or 0 to 1
  - Count number of gates for chip area

#### Obtaining design metrics of interest



# Implementation 2: Microcontroller and CCDPP



## Analysis of implementation 2

- Total execution time for processing one image:
  - 9.1 seconds
- Power consumption:
  - 0.033 watt
- Energy consumption:
  - 0.30 joule (9.1 s x 0.033 watt)
- Total chip area:
  - 98,000 gates

# Implementation 3: Microcontroller and CCDPP/Fixed-Point DCT

- 9.1 seconds still doesn't meet performance constraint of 1 second
- DCT opn prime candidate for improvement
  - Execution of implementation 2 shows microprocessor spends most cycles here
  - Could design custom hardware like we did for CCDPP
    - More complex so more design effort
  - Instead, will speed up DCT functionality by modifying behavior

# DCT floating-point cost



#### Floating-point cost

- DCT uses ~260 F.Pt. operations per pixel transformation
- 4096 (64 x 64) pixels per image
- 1 million floating-point operations per image
- No floating-point support with Intel 8051
  - Compiler must emulate
    - Generates procedures for each floating-point operation
       mult, add
    - Each procedure uses tens of integer operations
- Thus, > 10 million integer operations per image
- Procedures increase code size
- Fixed-point arithmetic can improve on this

# Fixed-point arithmetic



- Integer used to represent a real number
  - Constant number of integer's bits represents fractional portion of real number
    - More bits, more accurate the representation
  - Remaining bits represent portion of real number before decimal point

# Fixed-point arithmetic...



#### Translating a real constant to a fixed-point representation

- Multiply real value by 2 ^ (# of bits used for fractional part)
- Round to nearest integer
- E.g., represent 3.14 as 8-bit integer with 4 bits for fraction
  - $2^4 = 16$
  - $3.14 \times 16 = 50.24 \approx 50 = 00110010$
  - 16 (2<sup>4</sup>) possible values for fraction, each represents 0.0625 (1/16)
  - Last 4 bits (0010) = 2
  - $2 \times 0.0625 = 0.125$
  - 3(0011) + 0.125 = 3.125 ≈ 3.14 (more bits for fraction would increase accuracy)

# Fixed-point arithmetic operations



#### Addition

- Simply add integer representations
- E.g., 3.14 + 2.71 = 5.85
  - $3.14 \rightarrow 50 = 00110010$
  - $2.71 \rightarrow 43 = 00101011$
  - 50 + 43 = 93 = 01011101
  - $5(0101) + 13(1101) \times 0.0625 = 5.8125 \approx 5.85$

#### Multiply

- Multiply integer representations
- Shift result right by # of bits in fractional part
- E.g., 3.14 \* 2.71 = 8.5094
  - 50 \* 43 = 2150 = 100001100110
  - >> 4 = 10000110
  - $8(1000) + 6(0110) \times 0.0625 = 8.375 \approx 8.5094$
- Range of real values used limited by bit widths of possible resulting values

# Fixed-point implementation of CODEC



- COS\_TABLE gives 8-bit fixedpoint representation of cosine values
- 6 bits used for fractional portion
- Result of multiplications shifted right by 6

```
static unsigned char C(int h) { return h ? 64 : ONE_OVER_SQRT_TWO;}
static int F(int u, int v, short img[8][8]) {
   long s[8], r = 0;
   unsigned char x, j;
   for(x=0; x<8; x++) {
       s[x] = 0;
       for(j=0; j<8; j++)
            s[x] += (img[x][j] * COS_TABLE[j][v] ) >> 6;
   }
   for(x=0; x<8; x++) r += (s[x] * COS_TABLE[x][u]) >> 6;
   return (short)((((r * (((16*C(u)) >> 6) *C(v)) >> 6)) >> 6) >> 6);
}
```

```
static const char ONE_OVER_SQRT_TWO = 5;
static short xdata inBuffer[8][8], outBuffer[8][8], idx;
void CodecInitialize(void) { idx = 0; }
```

```
void CodecPushPixel(short p) {
   if( idx == 64 ) idx = 0;
   inBuffer[idx / 8][idx % 8] = p << 6; idx++;
}</pre>
```

```
void CodecDoFdct(void) {
   unsigned short x, y;
   for(x=0; x<8; x++)
        for(y=0; y<8; y++)
        outBuffer[x][y] = F(x, y, inBuffer);
   idx = 0;
}</pre>
```

# Implementation 3: Microcontroller and CCDPP/Fixed-Point DCT

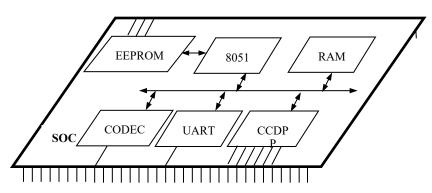


#### Analysis of implementation 3

- Use same analysis techniques as implementation 2
- Total execution time for processing one image:
  - 1.5 seconds
- Power consumption:
  - 0.033 watt (same as 2)
- Energy consumption:
  - 0.050 joule (1.5 s x 0.033 watt)
  - Battery life 6x longer!!
- Total chip area:
  - 90,000 gates
  - 8,000 less gates (less memory needed for code)

# Implementation 4: Microcontroller and CCDPP/DCT





- Performance close but not good enough
- Must resort to implementing CODEC in hardware
  - Single-purpose processor to perform DCT on 8 x 8 block

#### **Rewritten CODEC software**

# static unsigned char xdata C\_STAT\_REG \_at\_ 65527; static unsigned char xdata C\_CMND\_REG \_at\_ 65528; static unsigned char xdata C\_DATAI\_REG \_at\_ 65529; static unsigned char xdata C\_DATAO\_REG \_at\_ 65529; static unsigned char xdata C\_DATAI\_REG \_at\_ 65528; static unsigned char xdata C\_DATAI\_REG \_at\_ 65529; sta

## CODEC design

- 4 memory mapped registers
  - C\_DATAI\_REG/C\_DATAO\_REG
     used to push/pop 8 x 8 block into and out of CODEC
  - C\_CMND\_REG used to command CODEC
    - Writing 1 to this register invokes CODEC
  - C\_STAT\_REG indicates CODEC done and ready for next block
    - Polled in software
- Direct translation of C code to VHDL for actual hardware implementation
  - Fixed-point version used
- CODEC module in software changed similar to UART/CCDPP in implementation 2

# Implementation 4: Microcontroller and CCDPP/DCT



- Analysis of implementation 4
  - Total execution time for processing one image:
    - 0.099 seconds (well under 1 sec)
  - Power consumption:
    - 0.040 watt
    - Increase over 2 and 3 because SOC has another processor
  - Energy consumption:
    - 0.00040 joule (0.099 s x 0.040 watt)
    - Battery life 12x longer than previous implementation!!
  - Total chip area:
    - 128,000 gates, significant increase over previous implementations

# Summary of implementations



#### Implementation 3

- Close in performance
- Cheaper
- Less time to build

	lmpl 2	Impl 3	lmpl 4
Performance			
(second)	9.1	1.5	0.099
Power (watt)	0.033	0.033	0.040
Size (gate)	98,000	90,000	128,000
Energy (joule)	0.30	0.050	0.0040

#### Implementation 4

- Great performance and energy consumption
- More expensive and may miss time-to-market window
  - If DCT designed ourselves then increased NRE cost and time-to-market
  - If existing DCT purchased then increased IC cost

#### Which is better?

# Digital Camera -- Summary



#### Digital camera example

- Specifications in English and executable language
- Design metrics: performance, power and area

#### Several implementations

- Microcontroller: too slow
- Microcontroller and coprocessor: better, but still too slow
- Fixed-point arithmetic: almost fast enough
- Additional coprocessor for compression: fast enough, but expensive and hard to design
- Tradeoffs between hw/sw the main lesson of this course!

# Examples of Embedded Systems



#### We looked at details of

A simple Digital Camera

We will study microcontroller prog. with

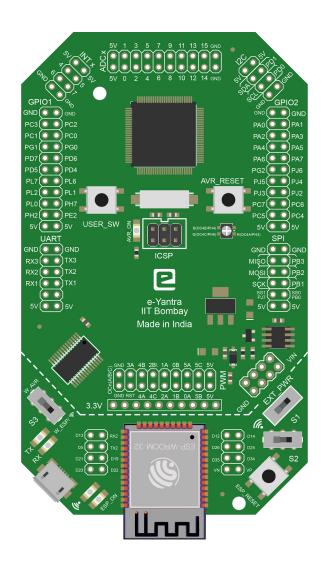
Atmega 2560
 Microcontroller & ESP32
 (to be studied in microcontroller workshop)

The world gets exciting...

 Apple iPad, intelligent transportation systems, service robots, ...







# eYFi-Mega Board

2 powerful controllers on single PCB
 [ATmega 2560 & ESP32]
 Wi-Fi compatible
 Bluetooth Low Energy
 Over-The-Air update
 12.5 W of Output Power
 On-board File Storage
 Compatible with FreeRTOS
 Arduino Programming Language compatible